

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
```


C:\Users\admin\anaconda3\lib\site-packages\scipy__init__.py:146: UserWarning: A NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy (detected version 1.23.5
warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")

```
In [3]: 1 cars = pd.read_csv('car.csv')
        2 cars.head()
```

Out[3]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	e
0	1	3	alfa-romero giulia	gas	std	two	convertible	rwd	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	rwd	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	rwd	
3	4	2	audi 100 ls	gas	std	four	sedan	fwd	
4	5	2	audi 100ls	gas	std	four	sedan	4wd	

5 rows × 26 columns



In [4]: 1 cars.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                205 non-null   int64
1   symboling              205 non-null   int64
2   CarName               205 non-null   object
3   fueltype              205 non-null   object
4   aspiration            205 non-null   object
5   doornumber            205 non-null   object
6   carbody               205 non-null   object
7   drivewheel            205 non-null   object
8   enginelocation        205 non-null   object
9   wheelbase             205 non-null   float64
10  carlength             205 non-null   float64
11  carwidth              205 non-null   float64
12  carheight             205 non-null   float64
13  curbweight            205 non-null   int64
14  enginetype            205 non-null   object
15  cylindernumber        205 non-null   object
16  enginesize            205 non-null   int64
17  fuelsystem            205 non-null   object
18  boreratio             205 non-null   float64
19  stroke                205 non-null   float64
20  compressionratio      205 non-null   float64
21  horsepower            205 non-null   int64
22  peakrpm               205 non-null   int64
23  citympg               205 non-null   int64
24  highwaympg            205 non-null   int64
25  price                 205 non-null   float64
dtypes: float64(8), int64(8), object(10)
memory usage: 41.8+ KB
```

```
In [5]: 1 cars.isnull().sum()
```


```
Out[5]: car_ID          0
        symboling      0
        CarName        0
        fueltype       0
        aspiration     0
        doornumber     0
        carbody        0
        drivewheel     0
        enginelocation 0
        wheelbase      0
        carlength      0
        carwidth       0
        carheight      0
        curbweight     0
        enginetype     0
        cylindernumber 0
        enginesize     0
        fuelsystem     0
        boreratio      0
        stroke         0
        compressionratio 0
        horsepower     0
        peakrpm        0
        citympg        0
        highwaympg     0
        price          0
        dtype: int64
```

```
In [6]: 1 plt.figure(figsize=(20,8))
2 plt.subplot(1,2,1)
3 plt.title("Car Price Distribution Plot")
4 sns.distplot(cars.price)
```

C:\Users\admin\anaconda3\lib\site-packages\seaborn\distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

warnings.warn(msg, FutureWarning)

```
Out[6]: <AxesSubplot:title={'center':'Car Price Distribution Plot'}, xlabel='price', ylabel='Density'>
```



In [7]: 1 cars.describe()

Out[7]:

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbweight
count	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000
mean	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	2555.565854
std	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	520.680204
min	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.000000
25%	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	2145.000000
50%	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	2414.000000
75%	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	2935.000000
max	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	4066.000000

In [8]: 1 cars.price.mean()

Out[8]: 13276.710570731706

In [9]: 1 cars.price.std()

Out[9]: 7988.85233174315

In [10]: 1 L = cars.price.mean() - 3*cars.price.std()
2 L

Out[10]: -10689.846424497744

In [11]: 1 U = cars.price.mean() + 3*cars.price.std()
2 U

Out[11]: 37243.267565961156

In [12]: 1 cars[(cars.price > U) | (cars.price < L)]

Out[12]:

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	drivewheel	eng
16	17	0	bmw x5	gas	std	two	sedan	rwd	
73	74	0	buick century special	gas	std	four	sedan	rwd	
74	75	1	buick regal sport coupe (turbo)	gas	std	two	hardtop	rwd	

3 rows × 26 columns

In [13]: 1 mcars = cars[(cars.price <= U) & (cars.price >= L)]

In [14]: 1 mcars.info()

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 202 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   car_ID                202 non-null   int64
1   symboling              202 non-null   int64
2   CarName               202 non-null   object
3   fueltype              202 non-null   object
4   aspiration            202 non-null   object
5   doornumber            202 non-null   object
6   carbody               202 non-null   object
7   drivewheel            202 non-null   object
8   enginelocation        202 non-null   object
9   wheelbase             202 non-null   float64
10  carlength             202 non-null   float64
11  carwidth              202 non-null   float64
12  carheight             202 non-null   float64
13  curbweight            202 non-null   int64
14  enginetype            202 non-null   object
15  cylindernumber        202 non-null   object
16  enginesize            202 non-null   int64
17  fuelsystem            202 non-null   object
18  boreratio             202 non-null   float64
19  stroke                202 non-null   float64
20  compressionratio      202 non-null   float64
21  horsepower            202 non-null   int64
22  peakrpm               202 non-null   int64
23  citympg               202 non-null   int64
24  highwaympg            202 non-null   int64
25  price                 202 non-null   float64
dtypes: float64(8), int64(8), object(10)
memory usage: 42.6+ KB
```

In [15]: 1 x = mcars.iloc[:, 20:24].values
2 y = mcars.iloc[:, 25].values

In [16]: 1 from sklearn.model_selection import train_test_split
2 x_train,x_test,y_train,y_test = train_test_split(x, y, test_size = 0.2,

In [17]: 1 from sklearn.linear_model import LinearRegression
2 model = LinearRegression()

In [18]: 1 model.fit(x_train,y_train)

Out[18]: LinearRegression()

```
In [19]: 1 ypred = model.predict(x_test)
          2 ypred
```

```
Out[19]: array([ 4175.45309753, 15843.56580946, 19611.01376363, 20380.85458482,
                20274.44466859,  9415.74662413, 14287.6930766 , 13548.51569155,
                16840.88783595,  7439.39262542, 10020.66612176, 25905.82412337,
                11009.80004235,  7332.25613048, 18042.99225923, 14287.6930766 ,
                6491.09930247, 14416.19472681,  8780.56313391,  6049.88783137,
                20558.82018667, 11033.06393096, 11440.65035548,  8003.82352924,
                10020.66612176, 25905.82412337,  9415.74662413, 10358.23807168,
                18576.32921943,  4175.45309753, 23581.08118075, 13784.50821037,
                6096.2578279 , 19611.01376363, 11190.83416232, 10085.02617381,
                14421.94723586, 14898.48545067, 10477.13737206, 13447.05590649,
                20201.65979893])
```

```
In [20]: 1 x_train.shape
```

```
Out[20]: (161, 4)
```

```
In [21]: 1 from sklearn.metrics import r2_score
          2 r2_score(y_test,ypred)
```

```
Out[21]: 0.7038832041380029
```

```
In [22]: 1 print(model.coef_)
          2 print(model.intercept_)
```

```
[ 575.25090463 112.5268254 -1.3594733 -335.78369938]
10877.103037647543
```

```
In [23]: 1 def scatter(x,fig):
          2     plt.subplot(5,2,fig)
          3     plt.scatter(cars[x], cars['price'])
          4     plt.title(x + ' vs Price')
          5     plt.ylabel('Price')
          6     plt.xlabel(x)
```

```
In [24]: 1 plt.figure(figsize = (10,20))
2 scatter('carlength', 1)
3 scatter('carwidth', 2)
4 scatter('carheight', 3)
5 scatter('curbweight', 4)
6 plt.tight_layout()
```

